

```
import pandas as pd
```

```
Patients_df = pd.read_csv('Patients.csv')
```

```
Patients_df
```

	Age	Diastolic	Gender	Height	LastName	Location	SelfAssessedHealthStatus
0	38	93	'Male'	71	'Smith'	'County General Hospital'	'Excellent'
1	43	77	'Male'	69	'Johnson'	'VA Hospital'	'Fair'
2	38	83	'Female'	64	'Williams'	'St. Mary's Medical Center'	'Good'
3	40	75	'Female'	67	'Jones'	'VA Hospital'	'Fair'
4	49	80	'Female'	64	'Brown'	'County General Hospital'	'Good'
...
95	25	99	'Male'	69	'Alexander'	'County General Hospital'	'Good'

```
!pip install plotly matplotlib seaborn --quiet
```


```
import plotly.express as px
import matplotlib
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
sns.set_style('darkgrid')
matplotlib.rcParams['font.size'] = 14
matplotlib.rcParams['figure.figsize'] = (10, 6)
matplotlib.rcParams['figure.facecolor'] = '#00000000'
```


```
Patients_df.Age.describe()
```

```
count    100.000000
mean      38.280000
std        7.215416
min       25.000000
25%       32.000000
50%       39.000000
75%       44.000000
max       50.000000
Name: Age, dtype: float64
```


```
fig = px.histogram(Patients_df,  
                   x='Age',  
                   marginal='box',  
                   nbins=26,  
                   title='Distribution of Age')  
fig.update_layout(bargap=0.1)  
fig.show()
```



```
fig = px.histogram(Patients_df,  
                  x='Weight',  
                  marginal='box',  
                  color_discrete_sequence=['red'],  
                  title='Distribution of Weight')  
fig.update_layout(bargap=0.1)  
fig.show()
```




```
fig = px.histogram(Patients_df,  
                   x='Systolic',  
                   marginal='box',  
                   color='Smoker',  
                   color_discrete_sequence=['green', 'grey'],  
                   title='Systolic Blood Pressure Level')  
fig.update_layout(bargap=0.1)  
fig.show()
```



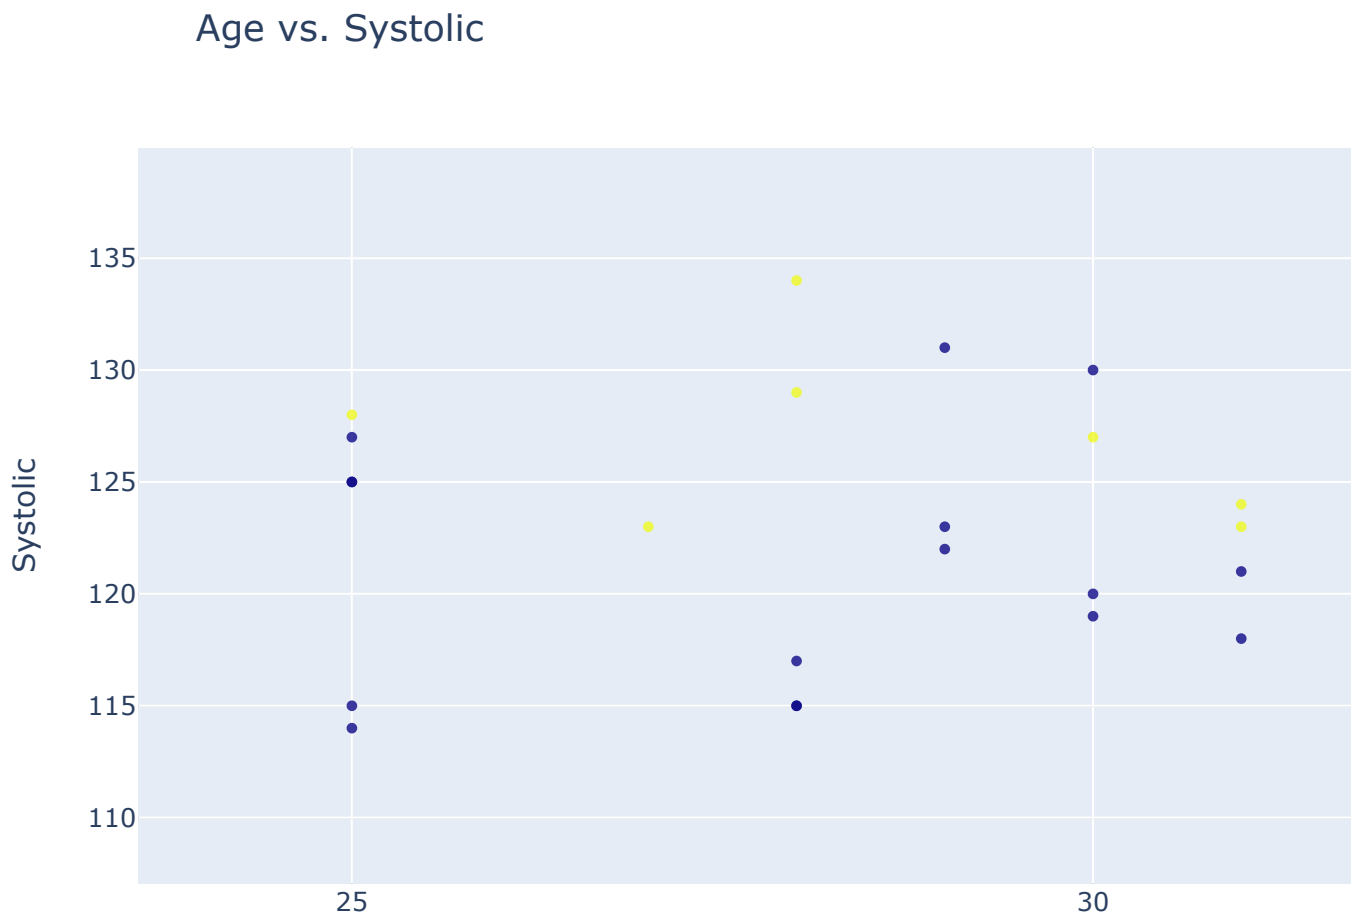
```
Patients_df.Smoker.value_counts()
```

```
0    66  
1    34  
Name: Smoker, dtype: int64
```

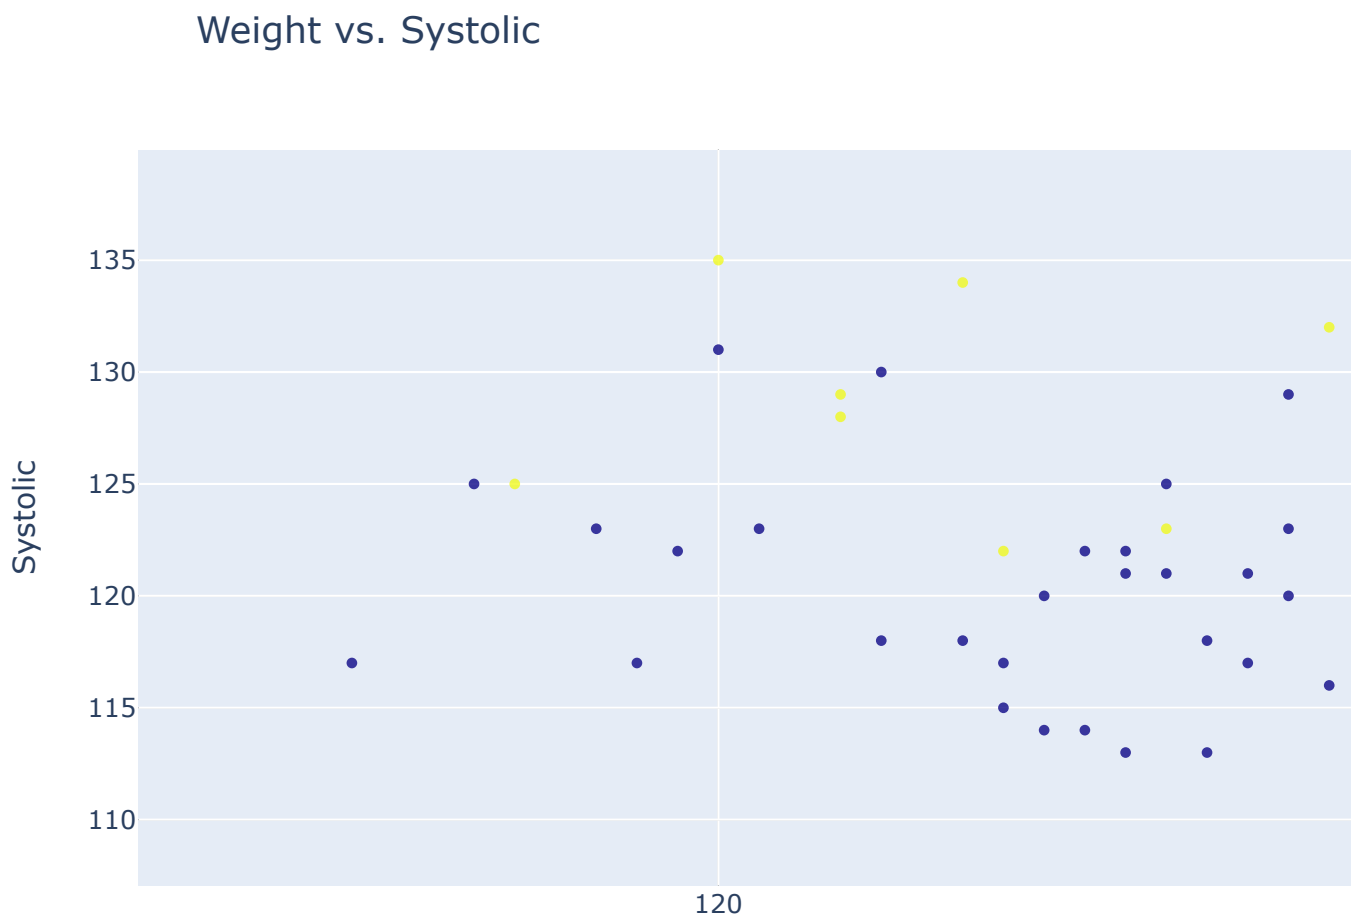
```
px.histogram(Patients_df, x='Smoker', color='Gender', title='Smoker')
```



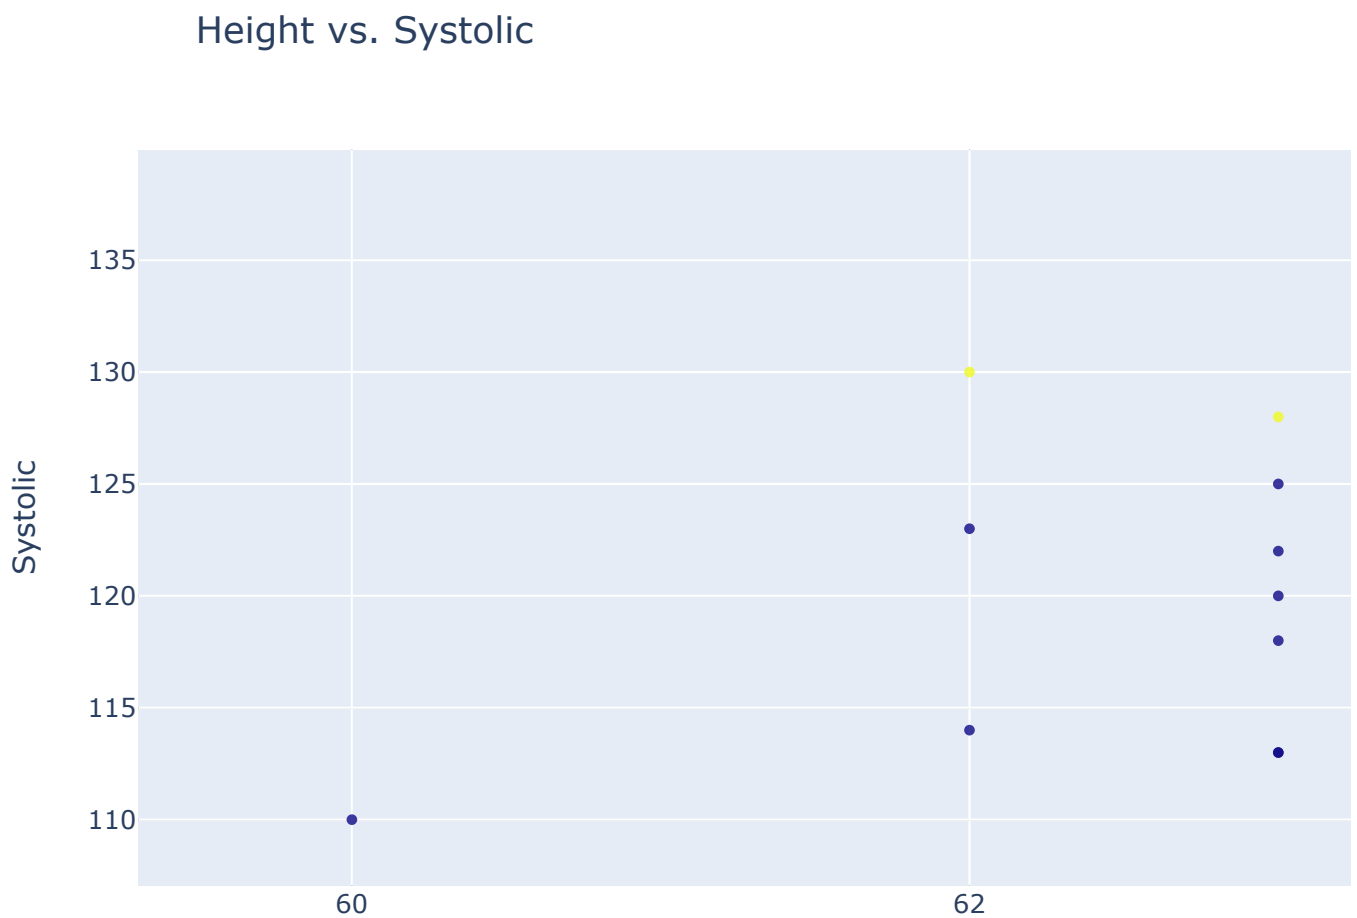
```
fig=px.scatter(Patients_df,·  
·····x='Age',·  
·····y='Systolic',·  
·····color='Smoker',·  
·····opacity=0.8,·  
·····hover_data=['Gender'],·  
·····title='Age vs. Systolic')  
fig.update_traces(marker_size=5)  
fig.show()
```



```
fig=px.scatter(Patients_df,·  
·····x='Weight',·  
·····y='Systolic',·  
·····color='Smoker',·  
·····opacity=0.8,·  
·····hover_data=['Gender'],·  
·····title='Weight vs. Systolic')  
fig.update_traces(marker_size=5)  
fig.show()
```




```
fig = px.scatter(Patients_df,  
                 x='Height',  
                 y='Systolic',  
                 color='Smoker',  
                 opacity=0.8,  
                 hover_data=['Gender'],  
                 title='Height vs. Systolic')  
fig.update_traces(marker_size=5)  
fig.show()
```



```
Patients_df.Systolic.corr(Patients_df.Age)
```

```
0.134126990051016
```

```
Patients_df.Systolic.corr(Patients_df.Weight)
```

```
0.15578811264461603
```

```
Patients_df.Systolic.corr(Patients_df.Height)
```

```
0.21407555251019653
```

```
smoker_values={'no':0, 'yes':1}
```

```
smoker_numeric=Patients_df.Smoker.map(smoker_values)
```


```
Patients_df.Systolic.corr(Patients_df.Smoker)
```

```
0.7063229808258187
```

```
Patients_df.corr()
```

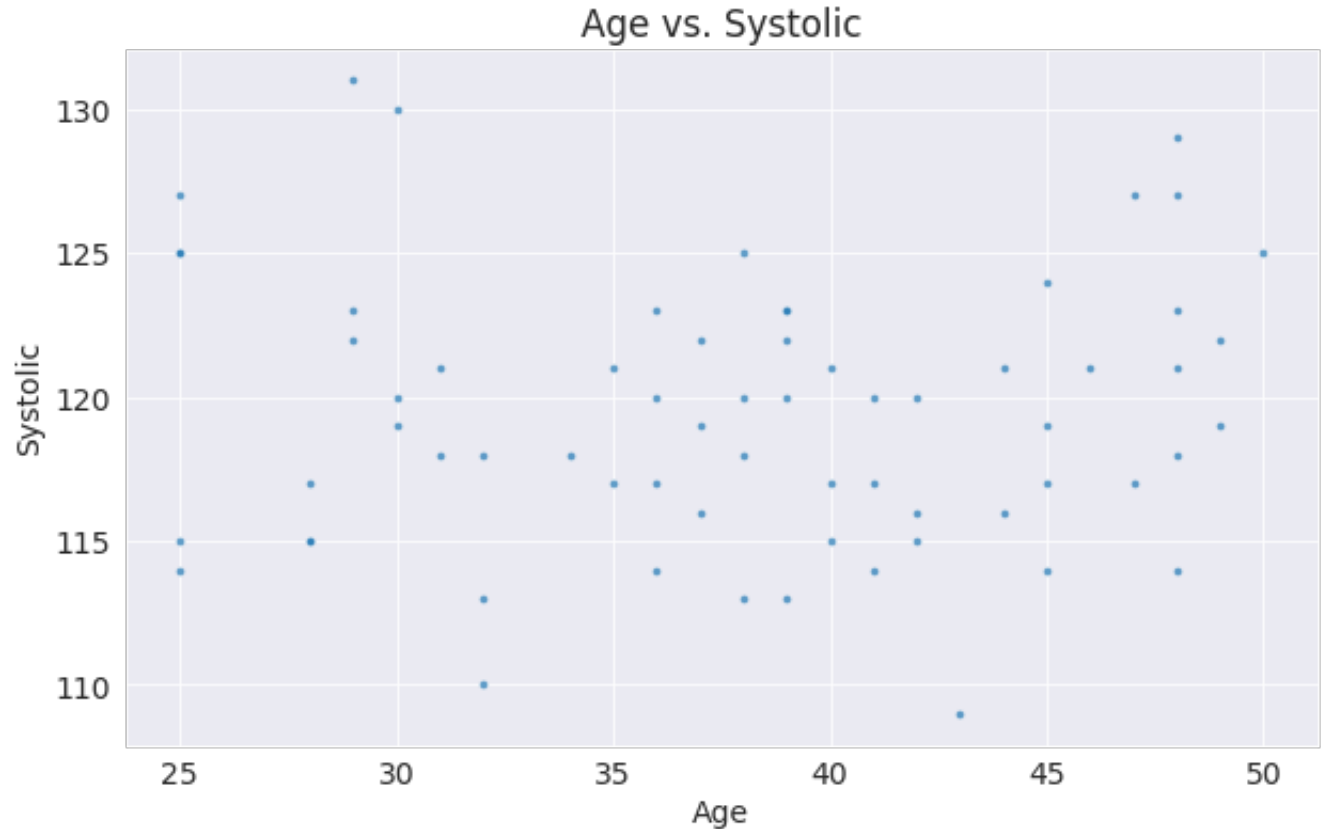
	Age	Diastolic	Height	Smoker	Systolic	Weight
Age	1.000000	0.080597	0.116002	0.060220	0.134127	0.091356
Diastolic	0.080597	1.000000	0.156819	0.723365	0.511843	0.222687
Height	0.116002	0.156819	1.000000	0.199114	0.214076	0.695970
Smoker	0.060220	0.723365	0.199114	1.000000	0.706323	0.215586
Systolic	0.134127	0.511843	0.214076	0.706323	1.000000	0.155788
Weight	0.091356	0.222687	0.695970	0.215586	0.155788	1.000000

```
sns.heatmap(Patients_df.corr(), cmap='Reds', annot=True)
plt.title('Correlation Matrix');
```



```
non_smoker_df = Patients_df[Patients_df.Smoker == 0]
```

```
plt.title('Age vs. Systolic')
sns.scatterplot(data=non_smoker_df, x='Age', y='Systolic', alpha=0.7, s=15);
```




```
def estimate_Systolic(Age, w, b):
    return w * Age + b
```

```
w = 2
b = 30
```

```
w = .3
b = .3
```


```
Ages = non_smoker_df.Age
estimated_Systolic = estimate_Systolic(Ages, w, b)
```

```
plt.plot(Ages, estimated_Systolic, 'r-o');  
plt.xlabel('Age');  
plt.ylabel('Estimated Systolic');
```




```
target = non_smoker_df.Systolic
```

```
plt.plot(Ages, estimated_Systolic, 'r', alpha=0.9);  
plt.scatter(Ages, target, s=8,alpha=0.8);  
plt.xlabel('Age');  
plt.ylabel('Systolic')  
plt.legend(['Estimate', 'Actual']);
```



```
def try_parameters(w, b):  
    Ages = non_smoker_df.Age  
    target = non_smoker_df.Systolic  
  
    estimated_Systolic = estimate_Systolic(Ages, w, b)  
  
    plt.plot(Ages, estimated_Systolic, 'r', alpha=0.9);  
    plt.scatter(Ages, target, s=8,alpha=0.8);  
    plt.xlabel('Age');  
    plt.ylabel('Systolic')  
    plt.legend(['Estimate', 'Actual']);
```

```
try_parameters(3, 1)
```



```
import numpy as np
```

```
def rmse(targets, predictions):  
    return np.sqrt(np.mean(np.square(targets - predictions)))
```

```
w = 3  
b = 3
```

```
targets = non_smoker_df['Systolic']  
predicted = estimate_Systolic(non_smoker_df.Age, w, b)
```


```
rmse(targets, predicted)
```

```
22.215882934185952
```

```
def try_parameters(w, b):  
    Ages = non_smoker_df.Age  
    target = non_smoker_df.Systolic  
    predictions = estimate_Systolic(Ages, w, b)  
  
    plt.plot(Ages, predictions, 'r', alpha=0.9);  
    plt.scatter(Ages, target, s=8, alpha=0.8);  
    plt.xlabel('Age');  
    plt.ylabel('Systolic')  
    plt.legend(['Prediction', 'Actual']);  
  
    loss = rmse(target, predictions)  
    print("RMSE Loss: ", loss)
```

```
try_parameters(3, 3)
```

```
RMSE Loss: 22.215882934185952
```



```
!pip install scikit-learn --quiet
```

```
from sklearn.linear_model import LinearRegression
```



```
model = LinearRegression()
```

```
help(model.fit)
```

Help on method fit in module sklearn.linear_model._base:

fit(X, y, sample_weight=None) method of sklearn.linear_model._base.LinearRegr
Fit linear model.

Parameters

X : {array-like, sparse matrix} of shape (n_samples, n_features)
Training data

y : array-like of shape (n_samples,) or (n_samples, n_targets)
Target values. Will be cast to X's dtype if necessary

sample_weight : array-like of shape (n_samples,), default=None
Individual weights for each sample

.. versionadded:: 0.17
parameter *sample_weight* support to LinearRegression.

Returns

self : returns an instance of self.

```
inputs = Patients_df[['Age']]
targets = Patients_df.Systolic
print('inputs.shape :', inputs.shape)
print('targes.shape :', targets.shape)
```

```
inputs.shape : (100, 1)
targes.shape : (100,)
```

```
model.fit(inputs, targets)
```

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
model.predict(np.array([[25],  
                        [36],  
                        [50]]))  
  
array([121.12285998, 122.49549102, 124.24247598])
```

```
predictions = model.predict(inputs)
```

```
predictions
```

```
array([122.7450603 , 123.3689835 , 122.7450603 , 122.99462958,  
      124.11769134, 123.74333742, 122.1211371 , 122.99462958,  
      121.4972139 , 121.87156782, 123.61855278, 123.24419886,  
      121.12285998, 122.86984494, 122.49549102, 123.9929067 ,  
      121.99635246, 121.37242926, 122.62027566, 124.24247598,  
      123.9929067 , 122.86984494, 123.11941422, 123.49376814,  
      121.4972139 , 121.12285998, 122.86984494, 121.12285998,  
      122.49549102, 121.74678318, 123.61855278, 122.99462958,  
      121.12285998, 123.86812206, 123.49376814, 123.9929067 ,  
      123.49376814, 122.37070638, 122.1211371 , 122.7450603 ,  
      122.86984494, 123.49376814, 123.49376814, 122.62027566,  
      123.61855278, 122.62027566, 121.74678318, 122.86984494,  
      123.24419886, 123.24419886, 124.11769134, 123.49376814,  
      123.3689835 , 123.86812206, 124.24247598, 122.7450603 ,  
      123.11941422, 123.61855278, 122.49549102, 122.7450603 ,  
      121.62199854, 121.4972139 , 121.74678318, 121.4972139 ,  
      121.62199854, 122.49549102, 123.61855278, 121.99635246,  
      121.87156782, 123.9929067 , 121.12285998, 122.99462958,  
      122.86984494, 123.11941422, 122.1211371 , 121.87156782,  
      122.37070638, 121.99635246, 123.24419886, 123.9929067 ,  
      122.24592174, 122.86984494, 121.4972139 , 121.62199854,  
      121.99635246, 122.86984494, 122.62027566, 124.11769134,  
      121.87156782, 122.62027566, 122.7450603 , 123.61855278,  
      121.74678318, 123.9929067 , 123.9929067 , 121.12285998,  
      123.49376814, 124.11769134, 123.61855278, 123.9929067 ])
```

```
rmse(targets, predictions)
```

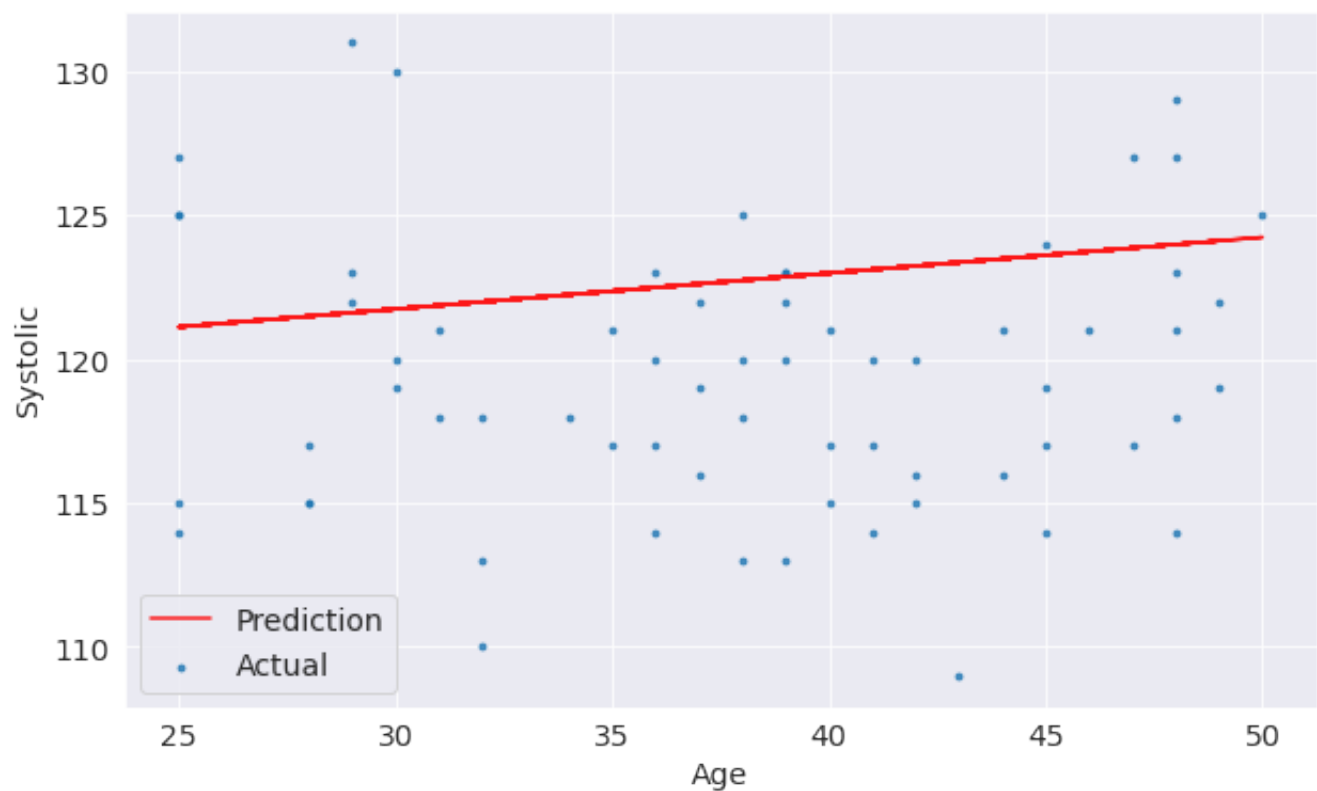
```
6.61883940804608
```

```
# w  
model.coef_  
  
array([0.12478464])
```

```
# b  
model.intercept_  
  
118.00324398156053
```

```
try_parameters(model.coef_, model.intercept_)
```

RMSE Loss: 5.775580444342182



```
# Create inputs and targets
inputs, targets = Patients_df[['Age', 'Weight']], Patients_df['Systolic']
```

```
# Create and train the model
model = LinearRegression().fit(inputs, targets)
```

```
# Generate predictions
predictions = model.predict(inputs)
```


```
# Compute loss to evaluate the model
loss = rmse(targets, predictions)
print('Loss:', loss)
```

Loss: 6.548450284122202

```
non_smoker_df.Systolic.corr(non_smoker_df.Weight)
```

-0.030177610486751022

```
fig = px.scatter(non_smoker_df, x='Weight', y='Systolic', title='Weight vs. Systol')
fig.update_traces(marker_size=5)
fig.show()
```




```
model.coef_, model.intercept_
```

(array([0.11248253, 0.03656693]), 112.84286158185228)

```
Patients_df.Systolic.corr(Patients_df.Smoker)
```

0.7063229808258187

```
fig = px.strip(Patients_df, x='Smoker', y='Systolic', title= "Smoker vs. Systolic")  
fig.update_traces(marker_size=4, marker_opacity=0.7)  
fig.show()
```



```
# Create inputs and targets
inputs, targets = Patients_df[['Age', 'Weight', 'Smoker']], Patients_df['Systolic']


# Create and train the model
model = LinearRegression().fit(inputs, targets)

# Generate predictions
predictions = model.predict(inputs)

# Compute loss to evaluate the model
loss = rmse(targets, predictions)
print('Loss:', loss)
```


Loss: 4.688172988059595

```
sns.barplot(data=Patients_df, x='Smoker', y='Systolic');
```



```
sns.barplot(data=Patients_df, x='Gender', y='Systolic')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fb7f56d1310>
```



Patients_df

	Age	Diastolic	Gender	Height	LastName	Location	SelfAssessedHealthStatus
0	38	93	'Male'	71	'Smith'	'County General Hospital'	'Excellent'
1	43	77	'Male'	69	'Johnson'	'VA Hospital'	'Fair'
2	38	83	'Female'	64	'Williams'	'St. Mary's Medical Center'	'Good'
3	40	75	'Female'	67	'Jones'	'VA Hospital'	'Fair'
4	49	80	'Female'	64	'Brown'	'County General Hospital'	'Good'
...
95	25	99	'Male'	69	'Alexander'	'County General Hospital'	'Good'

```
sex_codes = {'Female': 1, 'Male': 0}
```

```
Patients_df['sex_code'] = Patients_df.Gender.map(sex_codes)
```

Patients_df

	Age	Diastolic	Gender	Height	LastName	Location	SelfAssessedHealthStatus
0	38	93	'Male'	71	'Smith'	'County General Hospital'	'Excellent'
1	43	77	'Male'	69	'Johnson'	'VA Hospital'	'Fair'
2	38	83	'Female'	64	'Williams'	'St. Mary's Medical Center'	'Good'
3	40	75	'Female'	67	'Jones'	'VA Hospital'	'Fair'
4	49	80	'Female'	64	'Brown'	'County General Hospital'	'Good'
...
95	25	99	'Male'	69	'Alexander'	'County General Hospital'	'Good'

```
from sklearn import preprocessing
label = preprocessing.LabelEncoder()
```

```
Patients_df['Gender']= label.fit_transform(Patients_df['Gender'])
print(Patients_df['Gender'].unique())
```

```
[1 0]
```

Patients_df

	Age	Diastolic	Gender	Height	LastName	Location	SelfAssessedHealthStatus
0	38	93	1	71	'Smith'	'County General Hospital'	'Excellent'
1	43	77	1	69	'Johnson'	'VA Hospital'	'Fair'
2	38	83	0	64	'Williams'	'St. Mary's Medical Center'	'Good'
3	40	75	0	67	'Jones'	'VA Hospital'	'Fair'
4	49	80	0	64	'Brown'	'County General Hospital'	'Good'
...
95	25	99	1	69	'Alexander'	'County General Hospital'	'Good'

Patients_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 11 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   100 non-null    int64
1   Diastolic                             100 non-null    int64
2   Gender                               100 non-null    int64
3   Height                               100 non-null    int64
4   LastName                             100 non-null    object
5   Location                             100 non-null    object
6   SelfAssessedHealthStatus            100 non-null    object
7   Smoker                               100 non-null    int64
8   Systolic                             100 non-null    int64
9   Weight                               100 non-null    int64
10  sex_code                             0 non-null      float64
dtypes: float64(1), int64(7), object(3)
memory usage: 8.7+ KB
```

```
from sklearn.preprocessing import OneHotEncoder
from sklearn import preprocessing
Gender = ['Female', 'Male']
Location = ['County General Hospital', 'VA Hospital', "St. Mary's Medical Center"]
SelfAssessedHealthStatus = ['Fair', 'Good', 'Excellent']
enc = preprocessing.OneHotEncoder(categories=[Gender, Location, SelfAssessedHealthSt
X = [['Male', 'County General Hospital', 'Fair'], ['Female', 'VA Hospital', 'Good']
enc.fit(X)
OneHotEncoder(categories=[['Female', 'Male'],
                           ['County General Hospital', 'VA Hospital', "St. Mary's M
                           ['Fair', 'Good', 'Excellent']])
enc.transform([['Female', 'VA Hospital', 'Good']]).toarray()

array([[1., 0., 0., 1., 0., 0., 1., 0.]])
```

```
one_hot = enc.transform(Patients_df[['Location']]).toarray()
one_hot
```

[0., 0., 1.],
[0., 0., 1.],
[0., 0., 1.],
[0., 0., 1.],
[1., 0., 0.],
[0., 0., 1.],
[1., 0., 0.],
[1., 0., 0.],
[1., 0., 0.],
[0., 1., 0.],
[1., 0., 0.],
[0., 0., 1.],
[0., 0., 1.],
[0., 0., 1.],
[0., 1., 0.],
[0., 1., 0.],
[0., 0., 1.],
[0., 1., 0.],
[1., 0., 0.],
[0., 0., 1.],
[1., 0., 0.],
[1., 0., 0.],
[0., 0., 1.],
[0., 1., 0.],
[1., 0., 0.],
[1., 0., 0.],
[0., 0., 1.],
[1., 0., 0.],
[1., 0., 0.],
[0., 0., 1.],
[0., 1., 0.],
[0., 1., 0.],
[0., 1., 0.],
[0., 0., 1.],
[1., 0., 0.],
[0., 1., 0.],

```
[1., 0., 0.],  
[0., 0., 1.],  
[0., 1., 0.],  
[0., 0., 1.],  
[1., 0., 0.],  
[0., 1., 0.],  
[0., 0., 1.],  
[1., 0., 0.],  
[1., 0., 0.],  
[1., 0., 0.],  
[0., 0., 1.],  
[0., 1., 0.],  
[1., 0., 0.],  
[0., 0., 1.],  
[0., 1., 0.],  
[1., 0., 0.],  
[1., 0., 0.],  
[1., 0., 0.],  
[0., 0., 1.],  
[1., 0., 0.],  
[1., 0., 0.],  
[1., 0., 0.]])
```

```
Patients_df[['County·General·Hospital',·'VA·Hospital',·"St.·Mary's·Medical·Center"
```

Patients_df

	Age	Diastolic	Gender	Height	LastName	Location	SelfAssessedHealthStatus
0	38	93	1	71	'Smith'	'County General Hospital'	'Excellent'
1	43	77	1	69	'Johnson'	'VA Hospital'	'Fair'
2	38	83	0	64	'Williams'	'St. Mary's Medical Center'	'Good'
3	40	75	0	67	'Jones'	'VA Hospital'	'Fair'
4	49	80	0	64	'Brown'	'County General Hospital'	'Good'
...

```
from sklearn import preprocessing
enc = preprocessing.OneHotEncoder()
enc.fit(Patients_df[['SelfAssessedHealthStatus']])
enc.categories_
```

```
[array(['Excellent', 'Fair', 'Good', 'Poor'], dtype=object)]
```

```
one_hot = enc.transform(Patients_df[['SelfAssessedHealthStatus']]).toarray()
one_hot
```

```
[[1., 0., 0., 0.],
 [0., 0., 1., 0.],
 [1., 0., 0., 0.],
 [0., 0., 1., 0.],
 [0., 1., 0., 0.],
 [0., 0., 0., 1.],
 [0., 0., 1., 0.],
 [1., 0., 0., 0.],
 [0., 0., 1., 0.],
 [0., 0., 0., 1.],
 [0. 0. 1. 0.]
```

```
[0., 0., 1., 0.],  
[0., 0., 0., 1.],  
[1., 0., 0., 0.],  
[1., 0., 0., 0.],  
[1., 0., 0., 0.],  
[0., 0., 1., 0.],  
[0., 0., 1., 0.],  
[0., 0., 1., 0.],  
[0., 0., 1., 0.],  
[1., 0., 0., 0.],  
[0., 0., 1., 0.],  
[1., 0., 0., 0.],  
[0., 0., 1., 0.],  
[1., 0., 0., 0.],  
[0., 0., 1., 0.],  
[1., 0., 0., 0.],  
[0., 0., 1., 0.],  
[1., 0., 0., 0.],  
[1., 0., 0., 0.],  
[1., 0., 0., 0.],  
[1., 0., 0., 0.],  
[0., 0., 1., 0.],  
[0., 1., 0., 0.],  
[1., 0., 0., 0.],  
[0., 0., 0., 1.],  
[1., 0., 0., 0.],  
[0., 1., 0., 0.],  
[0., 1., 0., 0.],  
[0., 0., 0., 1.],  
[1., 0., 0., 0.],  
[0., 0., 1., 0.],  
[0., 0., 1., 0.],  
[1., 0., 0., 0.],  
[0., 0., 1., 0.],  
[0., 0., 1., 0.],  
[0., 0., 1., 0.],  
[1., 0., 0., 0.],  
[1., 0., 0., 0.],  
[0., 1., 0., 0.],  
[0., 0., 1., 0.],  
[1., 0., 0., 0.],  
[0., 0., 1., 0.],  
[0., 0., 1., 0.],  
[0., 0., 1., 0.],  
[1., 0., 0., 0.],  
[1., 0., 0., 0.],  
[0., 1., 0., 0.],  
[0., 0., 1., 0.],  
[1., 0., 0., 0.],  
[0., 0., 1., 0.],  
[0., 1., 0., 0.],  
[0., 0., 1., 0.],  
[0., 1., 0., 0.],  
[0., 0., 1., 0.],  
[0., 1., 0., 0.]])
```



```
Patients_df[['Excellent', 'Fair', 'Good', 'Poor']] = one_hot
```

```
Patients_df =Patients_df.dropna(axis=1)
```

```
Patients_df
```

	Age	Diastolic	Gender	Height	LastName	Location	SelfAssessedHealthStatu
0	38	93	1	71	'Smith'	'County General Hospital'	'Exceller
1	43	77	1	69	'Johnson'	'VA Hospital'	'Fa
2	38	83	0	64	'Williams'	'St. Mary's Medical Center'	'Goo
3	40	75	0	67	'Jones'	'VA Hospital'	'Fa
4	49	80	0	64	'Brown'	'County General Hospital'	'Goo
...

```
# Create inputs and targets

inputs, targets = Patients_df[['Age', 'Weight', 'Smoker', 'Gender', 'County·General

# Create and train the model
model = LinearRegression().fit(inputs, targets)

# Generate predictions
predictions = model.predict(inputs)

# Compute loss to evaluate the model
loss = rmse(targets, predictions)
print('Loss:', loss)
```

```
Loss: 4.530621990685625
```

```
Patients_df.isnull().sum()
```

```
Age                                0
Diastolic                         0
Gender                           0
Height                           0
LastName                         0
Location                         0
SelfAssessedHealthStatus         0
Smoker                           0
Systolic                         0
Weight                           0
County General Hospital          0
VA Hospital                      0
St. Mary's Medical Center        0
Excellent                        0
Fair                             0
Good                            0
Poor                             0
dtype: int64
```

```
model.coef_
```

```
array([ 0.09132731, -0.01515041,  9.87541468,  0.5488521 ,  0.57694361,
        0.17733968, -0.75428329,  0.55602622,  0.77978667, -2.11802701,
        0.78221412])
```



```
[ -1.15332416, -0.49171219, -0.0248029 ],
[ -1.43190487, -1.62643264, -0.37913001],
[ -1.29261451, -0.7564803 , 0.32952421],
[ -0.31758201, 1.32384052, 1.39250555],
[ 0.9360312 , -0.64300825, 1.03817844],
[ -0.87474344, -0.68083227, -2.50509268],
[ -1.0140338 , -0.90777636, -1.08778423],
[ 1.35390227, -0.64300825, -1.08778423],
[ -1.84977594, 1.21036847, -0.37913001],
[ 0.23957941, -1.0212484 , -1.08778423],
[ 0.10028906, 0.83212833, 1.74683266],
[ 0.37886977, -1.0212484 , -0.73345712],
[ -0.73545309, -1.47513658, -0.0248029 ],
[ -1.0140338 , 0.90777636, 1.74683266],
[ -0.45687237, -0.86995234, -1.08778423],
[ -0.87474344, 1.09689643, 0.32952421],
[ 0.51816013, 1.51296059, -0.37913001],
[ 1.35390227, -1.05907242, -1.08778423],
[ -0.59616273, 1.21036847, 0.32952421],
[ 0.10028906, 1.2860165 , 0.68385132],
[ -1.43190487, 1.32384052, 0.68385132],
[ -1.29261451, -1.2860165 , -1.08778423],
[ -0.87474344, -0.83212833, -1.44211134],
[ 0.10028906, 1.05907242, 0.32952421],
[ -0.17829166, -1.2860165 , -0.73345712],
[ 1.49319263, -1.17254446, -1.44211134],
[ -1.0140338 , -0.49171219, -0.37913001],
[ -0.17829166, -0.94560037, -0.73345712],
[ -0.0390013 , 1.13472045, 0.32952421],
[ 0.9360312 , 1.0212484 , 1.39250555],
[ -1.15332416, -1.13472045, 1.03817844],
[ 1.35390227, 0.7564803 , 1.39250555],
[ 1.35390227, -0.7564803 , -0.37913001],
[ -1.84977594, 0.64300825, 0.68385132],
[ 0.79674084, 1.2860165 , 0.68385132],
[ 1.49319263, 1.21036847, 1.03817844],
[ 0.9360312 , 0.68083227, 0.32952421],
[ 1.35390227, 0.86995234, -0.37913001]] )
```

```
cat_cols = ['Gender','Smoker','County General Hospital','VA Hospital',"St. Mary's I
categorical_data = Patients_df[cat_cols].values
```

```
inputs = np.concatenate((scaled_inputs, categorical_data), axis=1)
targets = Patients_df.Systolic
```

```
# Create and train the model
model = LinearRegression().fit(inputs, targets)
```

```
# Generate predictions
predictions = model.predict(inputs)
```

```
# Compute loss to evaluate the model
loss = rmse(targets, predictions)
print('Loss:', loss)
```

```
Loss: 4.445982271245141
```

```
Weights_df = pd.DataFrame({
    'feature': np.append(numeric_cols + cat_cols, 1),
    'weight': np.append(model.coef_, model.intercept_)
})
Weights_df.sort_values('weight', ascending=False)
```

	feature	weight
12	1	119.871389
4	Smoker	9.673087
2	Height	1.325387
9	Good	1.012690
11	Poor	0.885655
5	County General Hospital	0.863780
0	Age	0.576204
8	Excellent	0.426312
6	VA Hospital	0.007280
1	Weight	-0.354757
7	St. Mary's Medical Center	-0.871060
3	Gender	-1.479391
10	Fair	-2.324657

```
from sklearn.model_selection import train_test_split
```

```
inputs_train, inputs_test, targets_train, targets_test = train_test_split(inputs,
```

```
# Create and train the model
model = LinearRegression().fit(inputs_train, targets_train)

# Generate predictions
predictions_test = model.predict(inputs_test)

# Compute loss to evaluate the model
loss = rmse(targets_test, predictions_test)
print('Test Loss:', loss)
```

Test Loss: 4.20189185002708

```
# Generate predictions
predictions_train = model.predict(inputs_train)

# Compute loss to evaluate the model
loss = rmse(targets_train, predictions_train)
print('Training Loss:', loss)
```

Training Loss: 4.504253122353481

✓ 0s completed at 8:25 PM

