

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
plt.style.use('ggplot')
import warnings; warnings.simplefilter('ignore')
import seaborn as sns
import matplotlib
import matplotlib.pyplot as plt

sns.set_style('darkgrid')
matplotlib.rcParams['font.size'] = 14
matplotlib.rcParams['figure.figsize'] = (10, 6)
matplotlib.rcParams['figure.facecolor'] = '#00000000'

Patients_df = pd.read_csv('Patients.csv')

from sklearn import preprocessing
label = preprocessing.LabelEncoder()
.
Patients_df['Gender'] = label.fit_transform(Patients_df['Gender'])
print(Patients_df['Gender'].unique())

[1 0]

label = preprocessing.LabelEncoder()
.
Patients_df['Location'] = label.fit_transform(Patients_df['Location'])
print(Patients_df['Location'].unique())

[0 2 1]

label = preprocessing.LabelEncoder()

Patients_df['SelfAssessedHealthStatus'] = label.fit_transform(Patients_df['SelfAssessedHealthStatus'])
print(Patients_df['SelfAssessedHealthStatus'].unique())

[0 1 2 3]
```

```
from sklearn.preprocessing import OneHotEncoder
Gender = ['Female', 'Male']
Location = ['County General Hospital', 'VA Hospital', 'St. Mary's Medical Center']
SelfAssessedHealthStatus = ['Fair', 'Good', 'Excellent', 'Poor']
enc = preprocessing.OneHotEncoder(categories=[Gender, Location, SelfAssessedHealthSt
X = [['Male', 'County General Hospital', 'Fair'], ['Female', 'VA Hospital', 'Good']
enc.fit(X)
OneHotEncoder(categories=[['Female', 'Male'],
                           ['County General Hospital', 'VA Hospital', 'St. Mary's M
                           ['Fair', 'Good', 'Excellent']])
enc.transform([['Female', 'VA Hospital', 'Good']]).toarray()

array([[1., 0., 0., 1., 0., 0., 1., 0., 0.]])

del Patients_df['Diastolic']

del Patients_df['LastName']

enc = preprocessing.OneHotEncoder()
enc.fit(Patients_df[['Location']])
enc.categories_
one_hot = enc.transform(Patients_df[['Location']]).toarray()
one_hot
Patients_df[['County General Hospital', 'VA Hospital', 'St. Mary's Medical Center'

enc = preprocessing.OneHotEncoder()
enc.fit(Patients_df[['SelfAssessedHealthStatus']])
enc.categories_
one_hot = enc.transform(Patients_df[['SelfAssessedHealthStatus']]).toarray()
one_hot
Patients_df[['Excellent', 'Fair', 'Good', 'Poor']] = one_hot
```

Patients_df

	Age	Gender	Height	Location	SelfAssessedHealthStatus	Smoker	Systolic
0	38	1	71	0	0	1	124
1	43	1	69	2	1	0	109
2	38	0	64	1	2	0	125
3	40	0	67	2	1	0	117
4	49	0	64	0	2	0	122
...
95	25	1	69	0	2	1	128
96	44	1	69	2	2	1	124
97	49	1	70	0	1	0	119
98	45	1	68	0	2	1	136
99	48	1	66	0	1	0	114

100 rows x 8 columns

```
# Create inputs and targets
```

```
inputs, targets = Patients_df[['Age', 'Weight', 'Smoker', 'Gender', 'County General
```

```
# Create and train the model
```

```
model = LinearRegression().fit(inputs, targets)
```

```
# Generate predictions
```

```
predictions = model.predict(inputs)
```

```
# Compute loss to evaluate the model
```

```
loss = rmse(targets, predictions)
```

```
print('Loss:', loss)
```

```
from sklearn.preprocessing import StandardScaler
```

```
numeric_cols = ['Age', 'Weight', 'Height']
```

```

scaler = StandardScaler()
scaler.fit(Patients_df[numeric_cols])
scaled_inputs = scaler.transform(Patients_df[numeric_cols])
scaled_inputs

```

```

array([[ -0.0390013 ,  0.83212833,  1.39250555],
       [ 0.65745049,  0.34041613,  0.68385132],
       [-0.0390013 , -0.86995234, -1.08778423],
       [ 0.23957941, -0.79430431, -0.0248029 ],
       [ 1.49319263, -1.32384052, -1.08778423],
       [ 1.07532156, -0.45388818,  0.32952421],
       [-0.73545309, -0.45388818, -1.08778423],
       [ 0.23957941,  0.98342439,  0.32952421],
       [-1.43190487,  1.09689643,  0.32952421],
       [-1.0140338 , -0.83212833, -0.37913001],
       [ 0.9360312 , -0.98342439,  0.32952421],
       [ 0.51816013, -0.64300825, -0.37913001],
       [-1.84977594,  0.7564803 ,  1.39250555],
       [ 0.10028906,  1.81555271,  1.74683266],
       [-0.31758201, -0.94560037, -0.73345712],
       [ 1.35390227,  1.0212484 ,  1.39250555],
       [-0.87474344,  1.39948855,  0.68385132],
       [-1.57119523, -0.86995234,  0.68385132],
       [-0.17829166,  0.94560037,  1.03817844],
       [ 1.63248299,  0.68083227,  0.32952421],
       [ 1.35390227, -0.79430431, -0.73345712],
       [ 0.10028906, -1.39948855, -1.08778423],
       [ 0.37886977, -0.64300825, -1.79643845],
       [ 0.79674084, -0.30259212, -0.37913001],
       [-1.43190487, -1.17254446, -0.73345712],
       [-1.84977594,  1.32384052,  1.03817844],
       [ 0.10028906, -0.41606416, -1.44211134],
       [-1.84977594, -1.51296059, -1.44211134],
       [-0.31758201,  0.45388818,  0.32952421],
       [-1.15332416,  1.21036847, -0.0248029 ],
       [ 0.9360312 , -1.05907242,  1.03817844],
       [ 0.23957941, -0.64300825, -0.37913001],
       [-1.84977594, -0.60518424, -1.08778423],
       [ 1.21461191,  1.24819249,  1.03817844],
       [ 0.79674084,  1.47513658,  1.39250555],
       [ 1.35390227, -0.64300825, -0.37913001],
       [ 0.79674084,  1.43731256,  1.39250555],
       [-0.45687237, -1.36166453, -0.37913001],
       [-0.73545309,  0.98342439, -0.37913001],
       [-0.0390013 , -0.98342439, -1.44211134],
       [ 0.10028906,  0.37824015,  1.39250555],
       [ 0.79674084,  1.09689643,  0.68385132],
       [ 0.79674084,  0.56736022,  1.03817844],
       [-0.17829166,  1.51296059,  1.03817844],
       [ 0.9360312 ,  0.68083227, -0.0248029 ],
       [ 0.17829166,  0.71085600,  0.73345712]]

```

```
[ -0.17829100, -0.71805028, -0.73345712 ],
[ -1.15332416,  1.05907242,  0.32952421 ],
[  0.10028906, -1.24819249, -1.79643845 ],
[  0.51816013,  0.15129606,  1.03817844 ],
[  0.51816013,  0.94560037, -0.0248029 ],
[  1.49319263,  0.60518424,  0.32952421 ],
[  0.79674084, -0.68083227, -1.79643845 ],
[  0.65745049, -0.71865628, -1.08778423 ],
[  1.21461191, -0.2647681 , -0.37913001 ],
[  1.63248299,  1.21036847,  1.74683266 ],
[ -0.0390013 , -1.13472045, -1.44211134 ],
[  0.37886977, -0.7564803 , -0.37913001 ],
[  0.9360312 ,  0.60518424,  1.03817844 ],
[ -0.31758201,  0.98342439,  1.39250555 ],
[  0.00000000,  0.00000000,  0.00000000 ]
```

```
cat_cols = ['Gender', 'Smoker', 'County General Hospital', 'VA Hospital', "St. Mary's Hospital"]
categorical_data = Patients_df[cat_cols].values
```

```
from sklearn.model_selection import train_test_split
```

```
inputs_train, inputs_test, targets_train, targets_test = train_test_split(inputs,
```

```
# Create and train the model
model = LinearRegression().fit(inputs_train, targets_train)
```

```
# Generate predictions
predictions_test = model.predict(inputs_test)
```

```
# Compute loss to evaluate the model
loss = rmse(targets_test, predictions_test)
print('Test Loss:', loss)
```

```
# Generate predictions
predictions_train = model.predict(inputs_train)
```

```
# Compute loss to evaluate the model
loss = rmse(targets_train, predictions_train)
print('Training Loss:', loss)
```

Patients_df

	Age	Gender	Height	Location	SelfAssessedHealthStatus	Smoker	Systolic	
0	38	1	71	0	0	1	124	
1	43	1	69	2	1	0	109	
2	38	0	64	1	2	0	125	
3	40	0	67	2	1	0	117	
4	49	0	64	0	2	0	122	
...	
95	25	1	69	0	2	1	128	
96	44	1	69	2	2	1	124	
97	49	1	70	0	1	0	119	
98	45	1	68	0	2	1	136	
99	48	1	66	0	1	0	114	

100 rows × 8 columns

```
Patients_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 15 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                    100 non-null    int64
1   Gender                                100 non-null    int64
2   Height                                100 non-null    int64
3   Location                              100 non-null    int64
4   SelfAssessedHealthStatus             100 non-null    int64
5   Smoker                                100 non-null    int64
6   Systolic                              100 non-null    int64
7   Weight                                100 non-null    int64
8   County General Hospital               100 non-null    float64
9   VA Hospital                           100 non-null    float64
10  St. Mary's Medical Center             100 non-null    float64
11  Excellent                             100 non-null    float64
12  Fair                                  100 non-null    float64
13  Good                                  100 non-null    float64
14  Poor                                  100 non-null    float64
dtypes: float64(7), int64(8)
memory usage: 11.8 KB
```

```
data = {'Location' : ['County General Hospital', 'VA Hospital', "St. Mary's Medica
    ]}
```

```
df = pd.DataFrame(data,columns=['Location'])
```

```
print(df)
```

```
print(df.dtypes)
```

```
              Location
0   County General Hospital
1              VA Hospital
2  St. Mary's Medical Center
Location    object
dtype: object
```

```
'SelfAssessedHealthStatus' : ['Fair', 'Good', 'Excellent','Poor']
```

```

from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
Patients_df[['Age', 'Height',
             'Weight']] = scaler.fit_transform(Patients_df[['Age', 'Height',
             'Weight']])

```

```
Patients_df.head(5)
```

	Age	Gender	Height	Location	SelfAssessedHealthStatus	Smoker	Systolic
0	-0.039001	1	1.392506	0	0	1	
1	0.657450	1	0.683851	2	1	0	
2	-0.039001	0	-1.087784	1	2	0	
3	0.239579	0	-0.024803	2	1	0	
4	1.493193	0	-1.087784	0	2	0	

```

X = Patients_df.drop(['Systolic'], axis=1)
y = Patients_df['Systolic']

```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_s
```

```
from sklearn.linear_model import Lasso
```

```

reg = Lasso(alpha=0.5)
reg.fit(X_train, y_train)

```

```

Lasso(alpha=0.5, copy_X=True, fit_intercept=True, max_iter=1000,
      normalize=False, positive=False, precompute=False, random_state=None,
      selection='cyclic', tol=0.0001, warm_start=False)

```



```
print('Lasso Regression: R^2 score on training set', reg.score(X_train, y_train)*100)
print('Lasso Regression: R^2 score on test set', reg.score(X_test, y_test)*100)
```

```
Lasso Regression: R^2 score on training set 51.40426808603669
Lasso Regression: R^2 score on test set 35.038955672806814
```

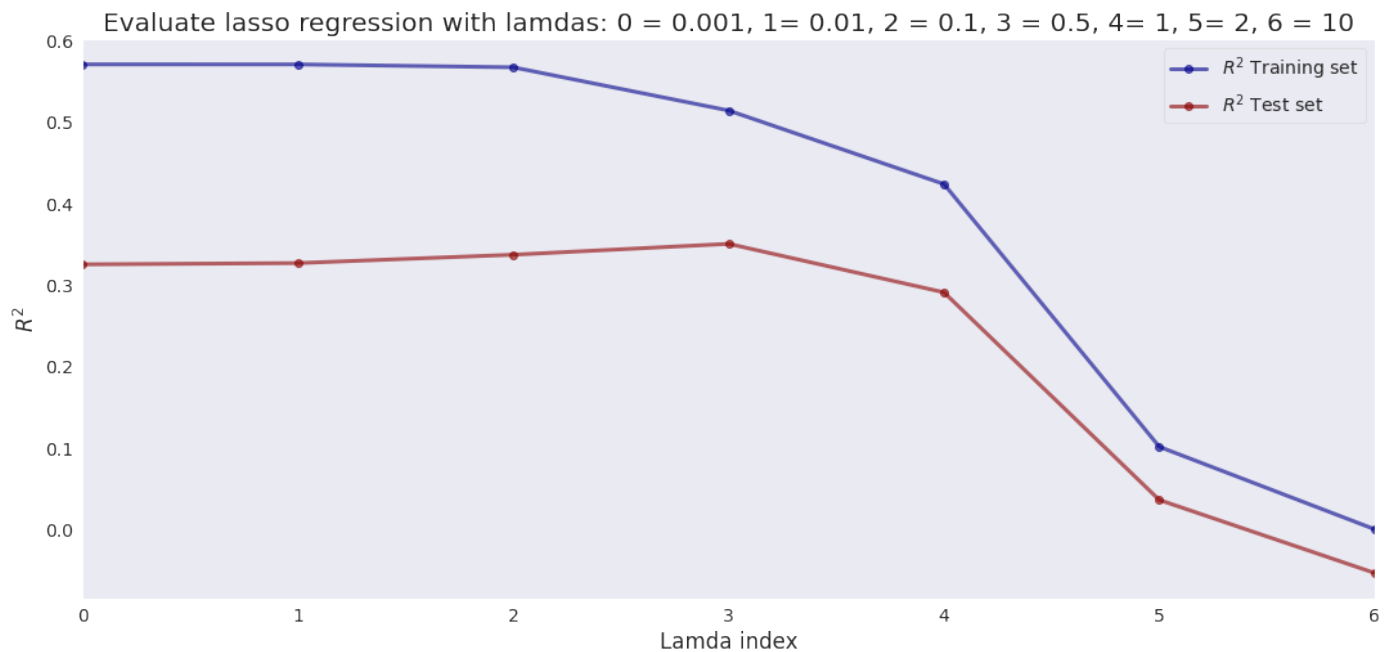
```
lambdas = (0.001, 0.01, 0.1, 0.5, 1, 2, 10)
l_num = 7
pred_num = X.shape[1]
```

```
# prepare data for enumerate
coeff_a = np.zeros((l_num, pred_num))
train_r_squared = np.zeros(l_num)
test_r_squared = np.zeros(l_num)
```

```
for ind, i in enumerate(lambdas):
    reg = Lasso(alpha = i)
    reg.fit(X_train, y_train)

    coeff_a[ind,:] = reg.coef_
    train_r_squared[ind] = reg.score(X_train, y_train)
    test_r_squared[ind] = reg.score(X_test, y_test)
```

```
plt.figure(figsize=(18, 8))
plt.plot(train_r_squared, 'bo-', label=r'$R^2$ Training set', color="darkblue", alpha=0.8)
plt.plot(test_r_squared, 'bo-', label=r'$R^2$ Test set', color="darkred", alpha=0.8)
plt.xlabel('Lamda index'); plt.ylabel(r'$R^2$')
plt.xlim(0, 6)
plt.title(r'Evaluate lasso regression with lamdas: 0 = 0.001, 1= 0.01, 2 = 0.1, 3 = 0.5, 4= 1, 5= 2, 6 = 10')
plt.legend(loc='best')
plt.grid()
```



```
df_lam = pd.DataFrame(test_r_squared*100, columns=['R_squared'])
df_lam['lambda'] = (lambdas)
# returns the index of the row where column has maximum value.
df_lam.loc[df_lam['R_squared'].idxmax()]
```

```
   R_squared    35.038956
   lambda      0.500000
Name: 3, dtype: float64
```

```
reg_best = Lasso(alpha = 0.1)
reg_best.fit(X_train, y_train)
reg_best.coef_
```

```
array([ 1.14804079, -0.          ,  1.04777052, -0.13619896,  0.45499808,
        8.75319922, -0.9764929 ])
```

```
from sklearn.metrics import mean_squared_error
mean_squared_error(y_test, reg_best.predict(X_test))
```

```
32.67545716126301
```

```
l_min = 0.05
l_max = 0.2
l_num = 20
lambdas = np.linspace(l_min, l_max, l_num)
```

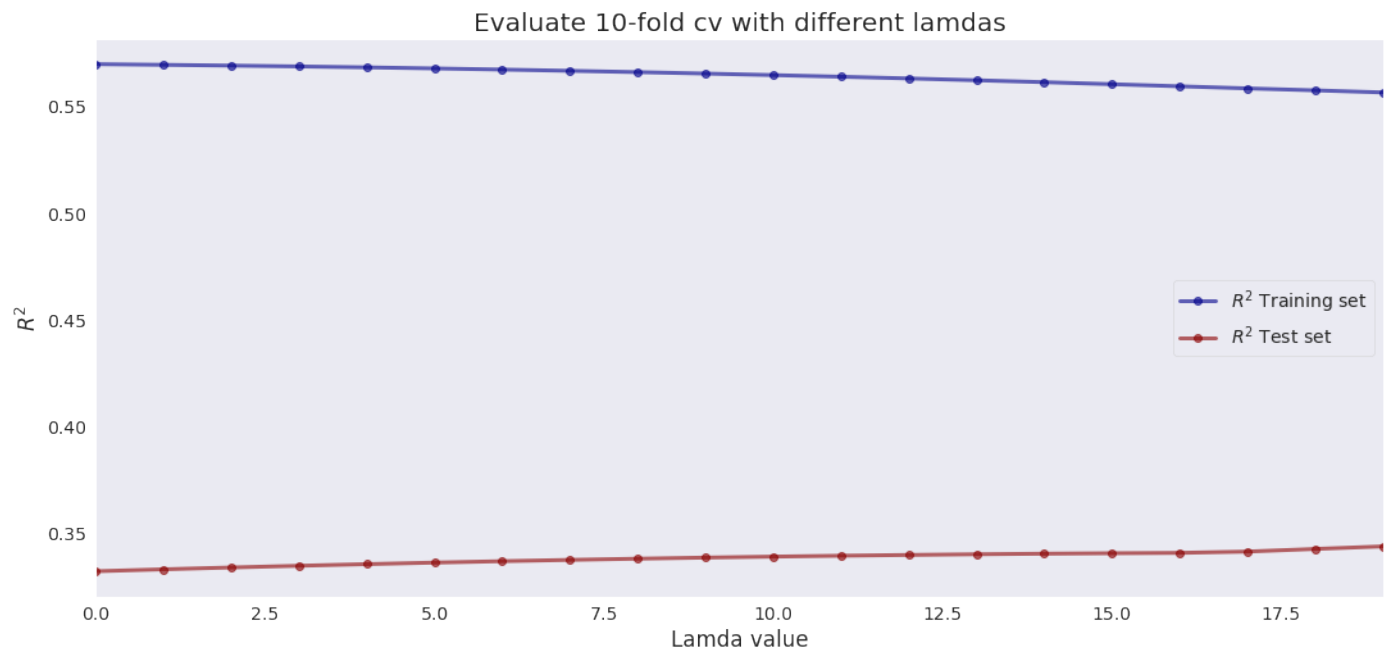
```
train_r_squared = np.zeros(l_num)
test_r_squared = np.zeros(l_num)
```

```
pred_num = X.shape[1]
coeff_a = np.zeros((l_num, pred_num))
```

```
from sklearn.model_selection import cross_val_score
```

```
for ind, i in enumerate(lambdas):  
    reg = Lasso(alpha = i)  
    reg.fit(X_train, y_train)  
    results = cross_val_score(reg, X, y, cv=10, scoring="r2")  
  
    train_r_squared[ind] = reg.score(X_train, y_train)  
    test_r_squared[ind] = reg.score(X_test, y_test)
```

```
# Plotting
plt.figure(figsize=(18, 8))
plt.plot(train_r_squared, 'bo-', label=r'$R^2$ Training set', color="darkblue", alpha=0.8)
plt.plot(test_r_squared, 'bo-', label=r'$R^2$ Test set', color="darkred", alpha=0.8)
plt.xlabel('Lamda value'); plt.ylabel(r'$R^2$')
plt.xlim(0, 19)
plt.title(r'Evaluate 10-fold cv with different lamdas')
plt.legend(loc='best')
plt.grid()
```



```
df_lam = pd.DataFrame(test_r_squared*100, columns=['R_squared'])
df_lam['lambda'] = (lambdas)
# returns the index of the row where column has maximum value.
df_lam.loc[df_lam['R_squared'].idxmax()]
```

```
R_squared    34.376749
lambda        0.200000
Name: 19, dtype: float64
```

```
# Best Model
reg_best = Lasso(alpha = 0.144737)
reg_best.fit(X_train, y_train)
```

```
Lasso(alpha=0.144737, copy_X=True, fit_intercept=True, max_iter=1000,
      normalize=False, positive=False, precompute=False, random_state=None,
      selection='cyclic', tol=0.0001, warm_start=False)
```

```
from sklearn.metrics import mean_squared_error
```

```
mean_squared_error(y_test, reg_best.predict(X_test))
```

```
32.54105584096965
```

```
reg_best.coef_
```

```
array([ 1.11565225, -0.          ,  0.87507912, -0.06084871,  0.41829045,
        8.571702   , -0.78587678])
```

✓ 0s completed at 1:21 PM

